

LLM Reasoning

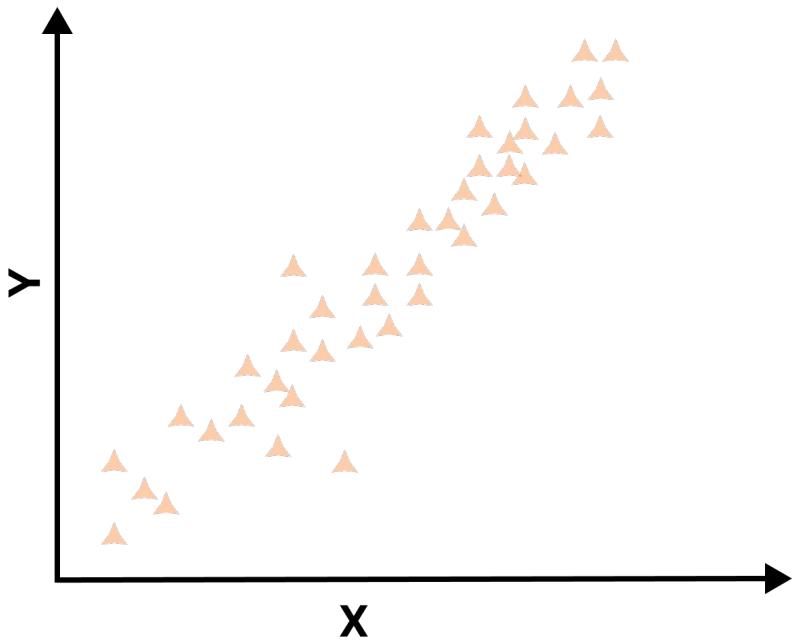
Course 1: Sampling Methods

Yaya Sy

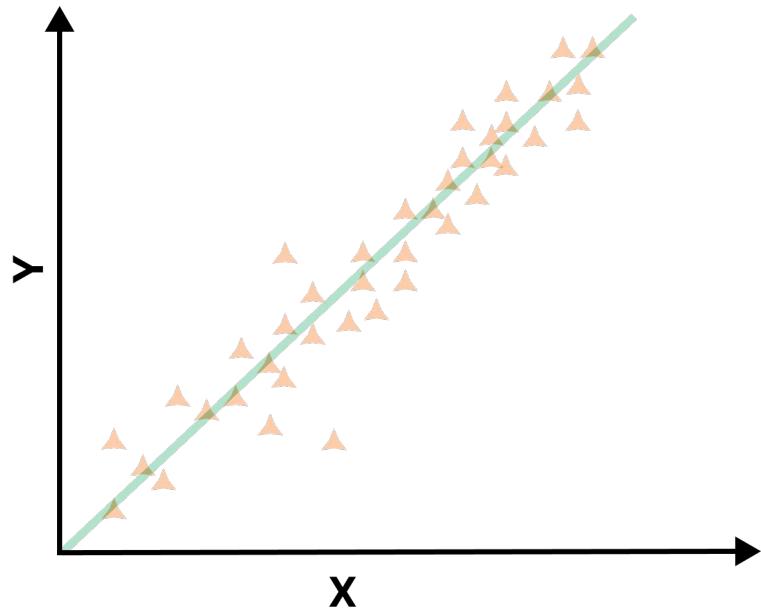
12/2/2025

Introduction

LLM = function that predicts the next token

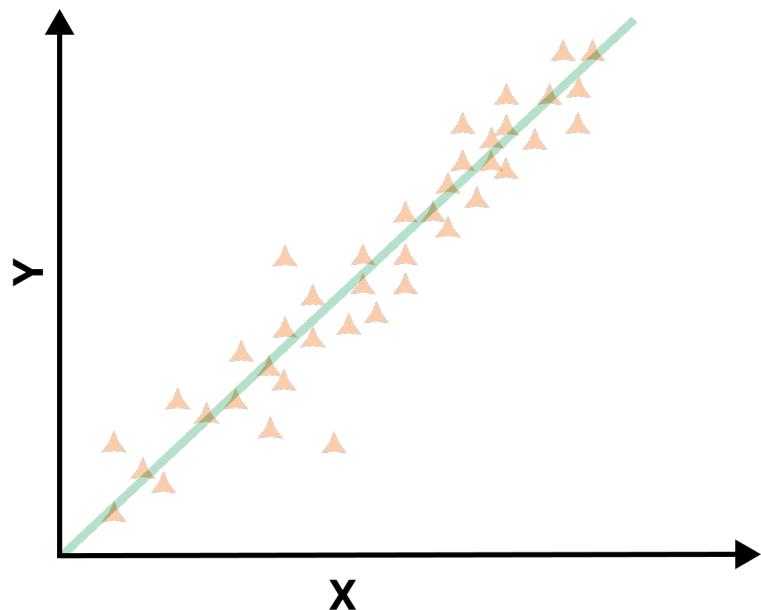


LLM = function that predicts the next token



$$f(x) = ax + b$$

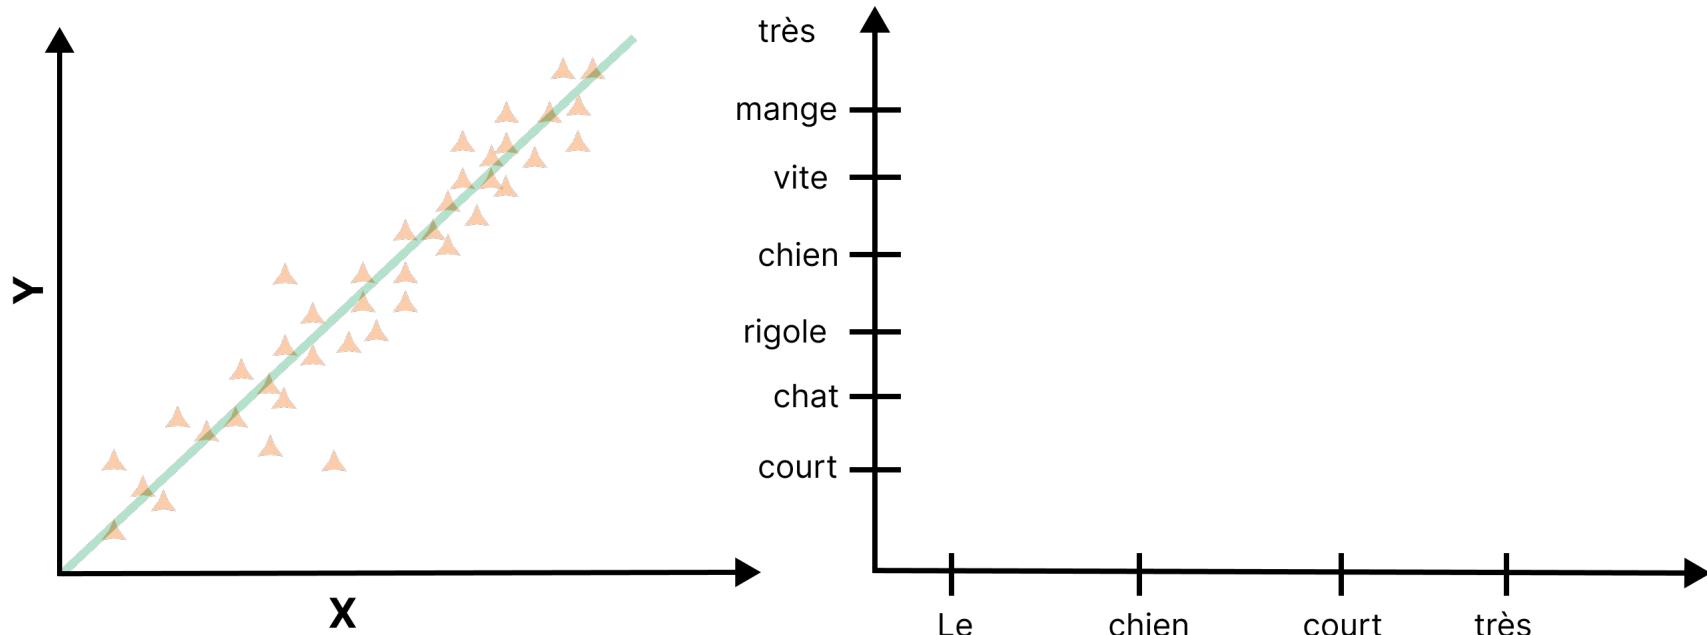
LLM = function that predicts the next token



$$f(x) = ax + b$$

2 parameters a, b

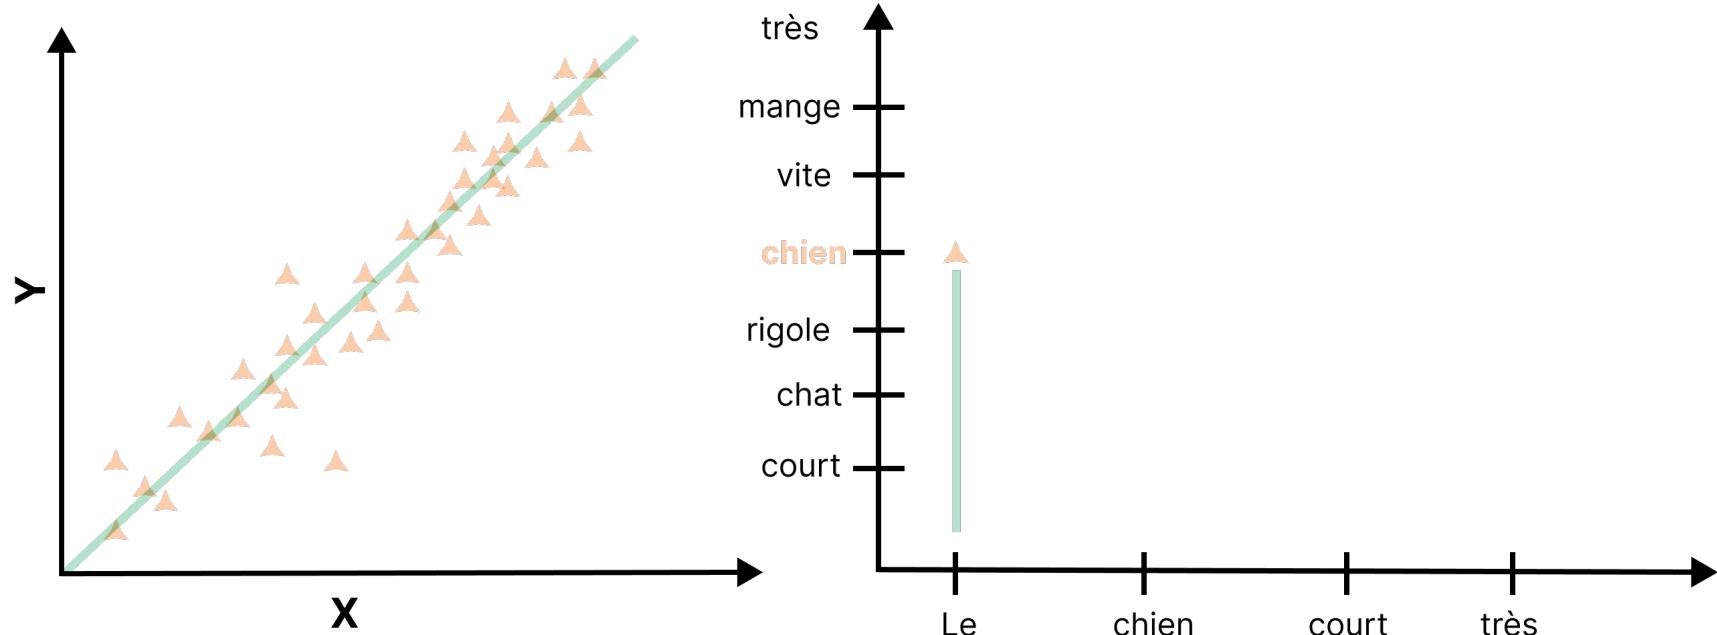
LLM = function that predicts the next token



$$f(x) = ax + b$$

2 parameters a, b

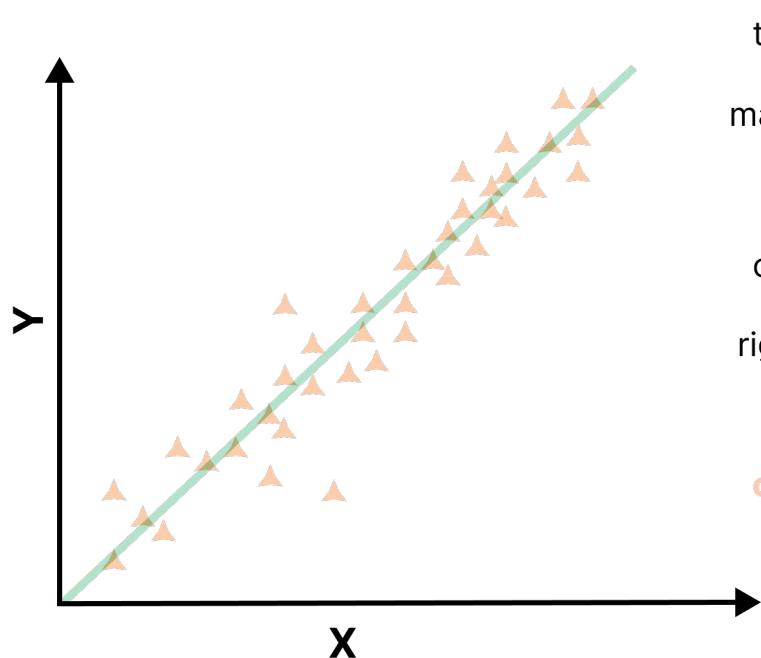
LLM = function that predicts the next token



$$f(x) = ax + b$$

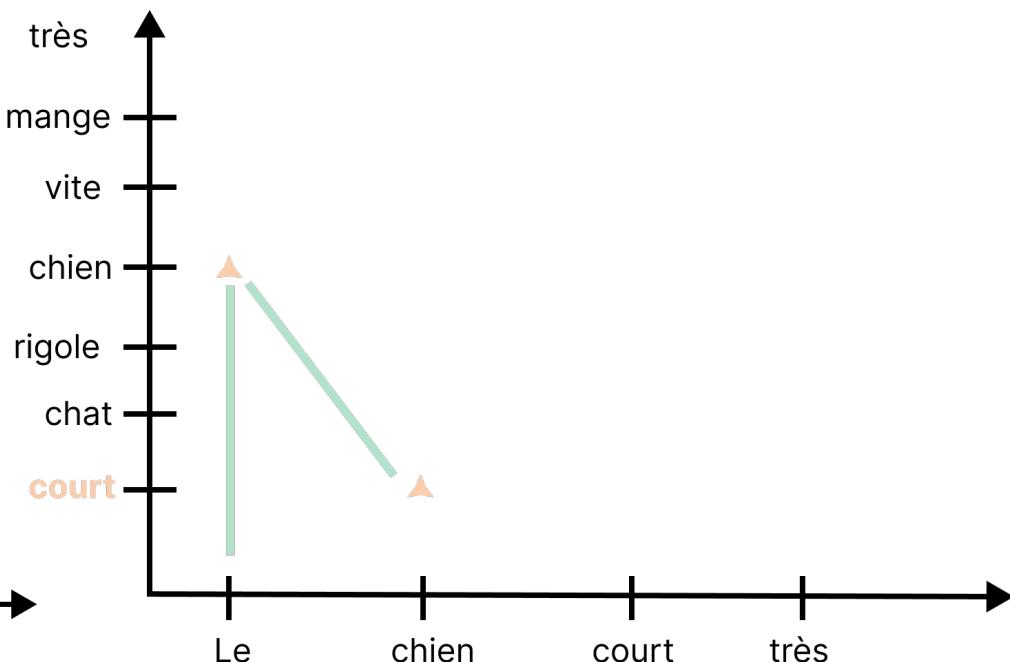
2 parameters a, b

LLM = function that predicts the next token

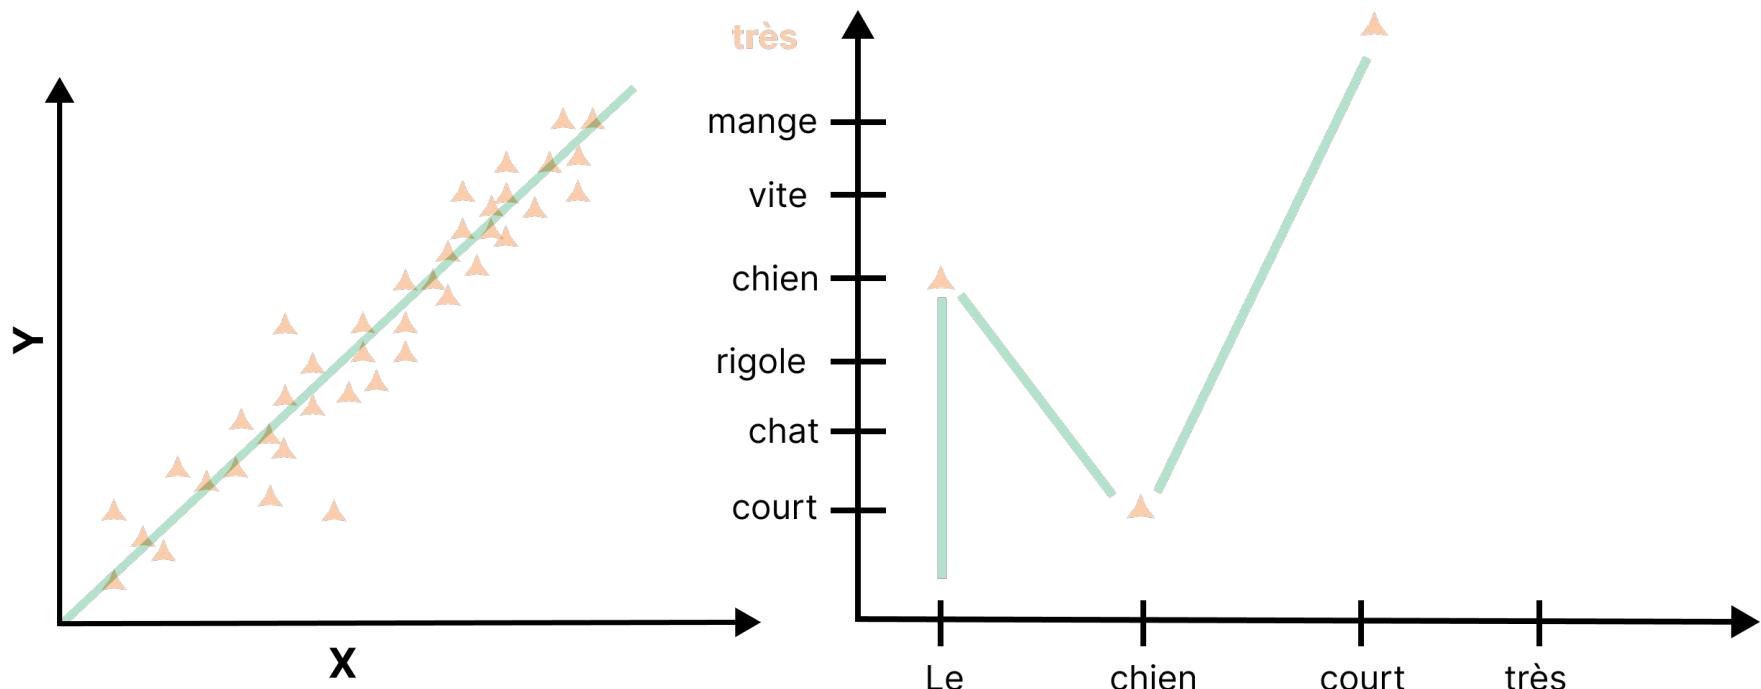


$$f(x) = ax + b$$

2 parameters a, b



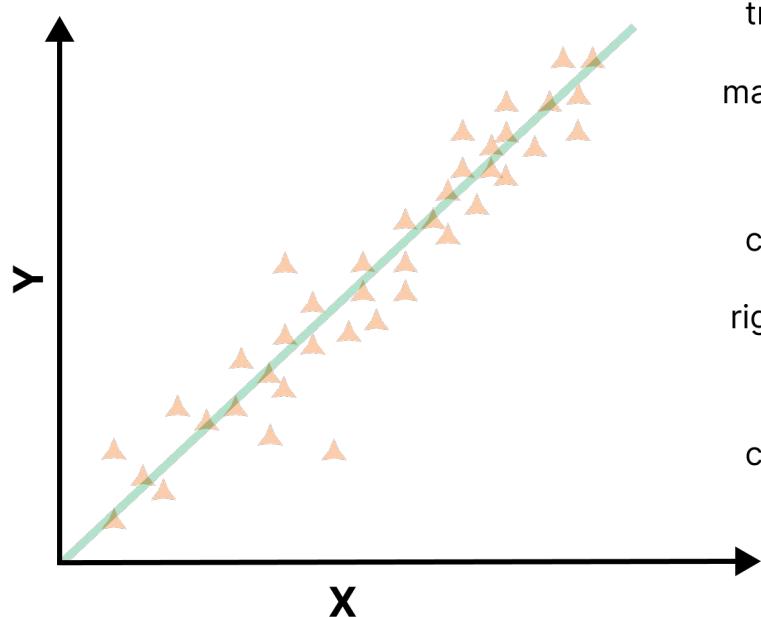
LLM = function that predicts the next token



$$f(x) = ax + b$$

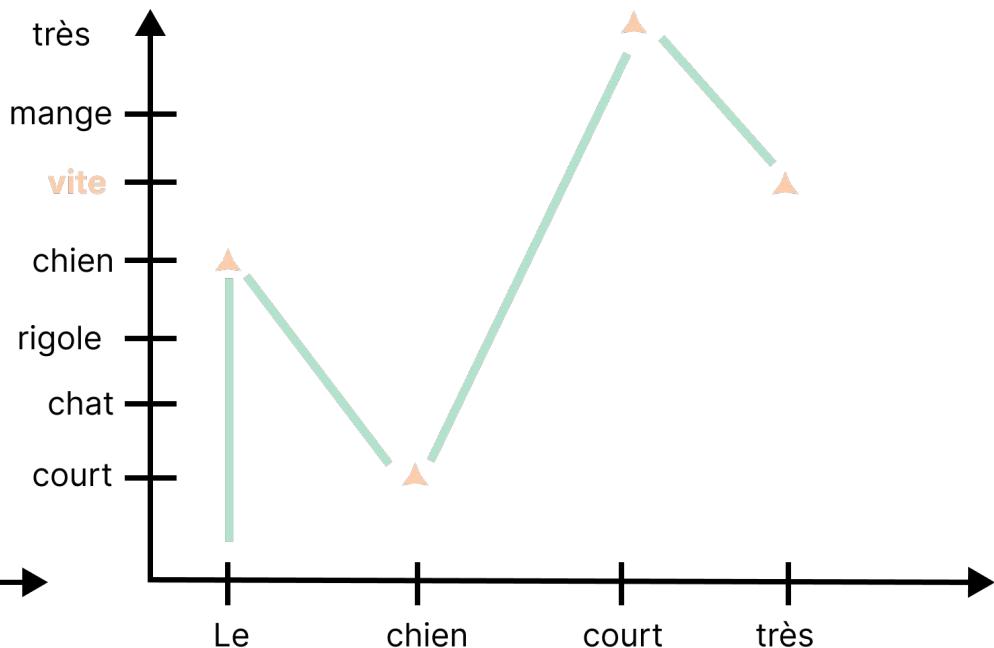
2 parameters a, b

LLM = function that predicts the next token

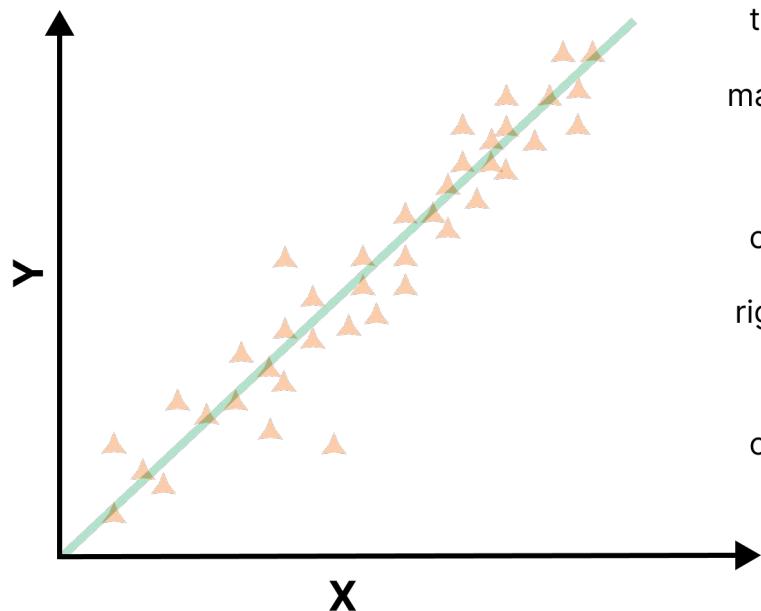


$$f(x) = ax + b$$

2 parameters a, b

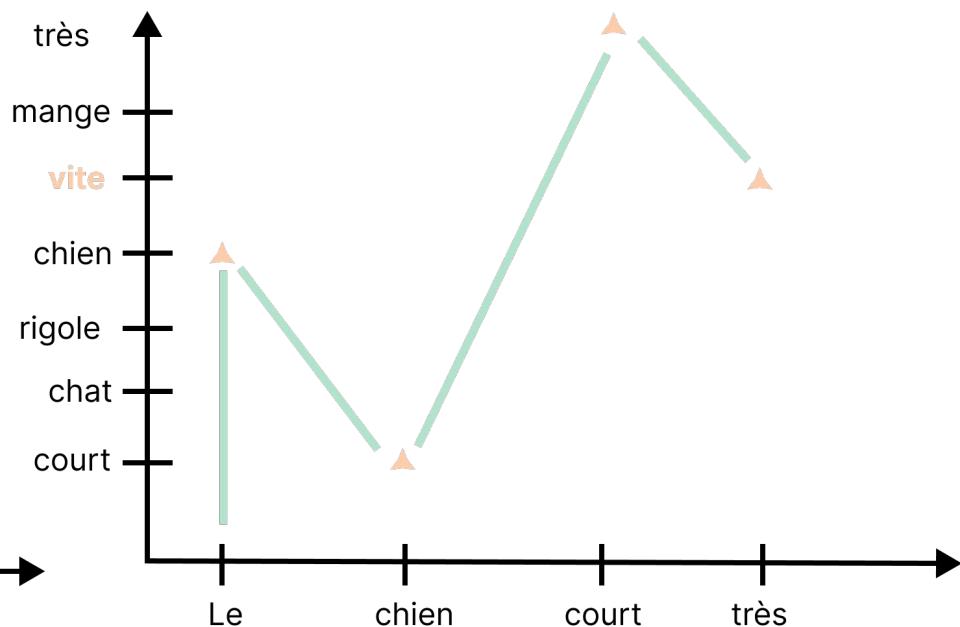


LLM = function that predicts the next token



$$f(x) = ax + b$$

2 parameters: a, b



$$f(\mathbf{x}) = \text{LLM}(\mathbf{x})$$

billion of parameters!

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court**

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite**

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite**
- Classification: Super le film. J'ai adoré ! →

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite.**
- Classification: Super le film. J'ai adoré ! → **Positive**

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite.**
- Classification: Super le film. J'ai adoré ! → **Positive**
- Chat: Hello, who are you? →

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite.**
- Classification: Super le film. J'ai adoré ! → **Positive**
- Chat: Hello, who are you? → **I**

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite.**
- Classification: Super le film. J'ai adoré ! → **Positive**
- Chat: Hello, who are you? → **I**

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite.**
- Classification: Super le film. J'ai adoré ! → **Positif**
- Chat: Hello, who are you? → **I am**

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite.**
- Classification: Super le film. J'ai adoré ! → **Positif**
- Chat: Hello, who are you? → **I am an**

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite.**
- Classification: Super le film. J'ai adoré ! → **Positif**
- Chat: Hello, who are you? → **I am an AI**

Language Models are Unsupervised Multitask Learners

Alec Radford ^{*} ¹ Jeffrey Wu ^{*} ¹ Rewon Child ¹ David Luan ¹ Dario Amodei ^{**} ¹ Ilya Sutskever ^{**}

Many natural language tasks can be solved with next-token prediction.

- Translation: The dog runs fast → **Le chien court vite.**
- Classification: Super le film. J'ai adoré ! → **Positif**
- Chat: Hello, who are you? → **I am an AI assistant.**

Learning to predict is learning a probability function

An LLM as a probability function

An LLM is a function that predicts the next token t given its parameters θ and the context c :

$$LLM(t|c, \theta)$$

And this function is a probability function:

$$p(t|c, \theta)$$

θ are the neural network parameters which is the Transformer.

c is the context vector computed by the *Transformer*, (e.g, $c = Transformer(text)$)

$$p(t|c, \theta) = \frac{score(c, t)}{\sum_{t' \in \mathcal{V}} score(c, t')}$$

An LLM as a probability function

$$p(t|c, \theta) = \frac{\text{score}(c, t)}{\sum_{t' \in \mathcal{V}} \text{score}(c, t')}$$

$\text{score}(c, t)$ measures how strong the context c is associated with the token t

$\text{score}(c, t')$ measures how strong the context c is associated with any other token t' in the vocabulary \mathcal{V}

$\text{score}(c, t) = \mathbf{c} \cdot \mathbf{w}_t^T$ The score can be defined as a dot product between the context vector and the token weight. This is also called ***logits***

An LLM as a probability function

$$p(t|c, \theta) = \frac{c \cdot w_t^T}{\sum_{t' \in \mathcal{V}} c \cdot w_{t'}^T}$$

- A probability is between 0 and 1, and cannot be negative. We need positive scores.
- *The dot product score returns any real number*, it can be positive or negative. It is not suited for a probability function
- To always have positive scores, the solution is to take the *exponential of the scores*:

$$p(t|c, \theta) = \frac{\exp(c \cdot w_t^T)}{\sum_{t' \in \mathcal{V}} \exp(c \cdot w_{t'}^T)}$$

- This is **softmax**

An LLM as a probability function

- Now that we know that an LLM is a probability function, we can use it to compute the probability of any text.
- The probability of the text “ I love this <e>” is:

$$\begin{aligned} p(<\text{b}> \text{ I love this } <\text{e}>) &= p(t = \text{I} \mid c = <\text{b}>, \theta) \\ &\quad \times p(t = \text{love} \mid c = <\text{b}> \text{ I}, \theta) \\ &\quad \times p(t = \text{this} \mid c = <\text{b}> \text{ I love}, \theta) \\ &\quad \times p(t = <\text{e}> \mid c = <\text{b}> \text{ I love this}, \theta) \end{aligned}$$

An LLM as a probability function

→ In general, the probability of any text s under the LLM θ :

$$p(s) = \prod_{i=1}^N p(t = s_i \mid c = s_{<i}, \theta)$$

→ In practice, for numerical stability reason, we compute the log-probability:

$$\log p(s) = \sum_{i=1}^N \log p(t = s_i \mid c = s_{<i}, \theta)$$

How this probability function is learned?

How an LLM learns a good probability function ?

To summarize:

An LLM is a probability function p based on the Transformer architecture with parameters θ , that computes the probability of any sequence $p(s)$

But we don't know in advance what is the probability function governing the texts. The LLMs needs to learn one.

How an LLM learns a good probability function?

The LLM is trained to learn the probability function that maximizes the probability of the training data \mathcal{D} . We call this approach ***Maximum Likelihood Estimation (MLE)***.

$$\log P(\mathcal{D}; \theta) = \sum_{s \in \mathcal{D}} \sum_{i=1}^{|s|} \log p(t = s_i \mid c = s_{<i}, \theta)$$

Then the LLM θ learns to maximize the probability of the training data \mathcal{D}

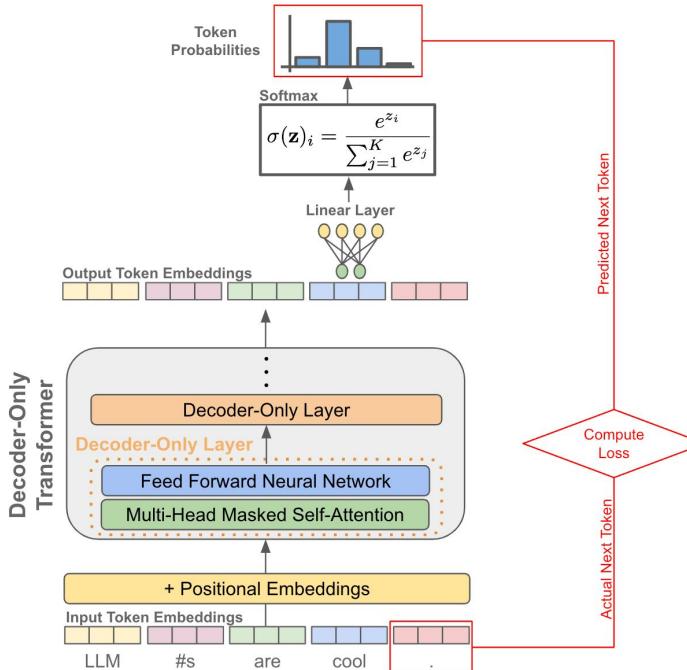
$$\hat{\theta} = \arg \max_{\theta} \log P(\mathcal{D}; \theta)$$

Maximizing the log-probabilities is equivalent to minimizing the negative log-probabilities, which is equivalent to ***cross-entropy loss***

The different steps of LLM Training

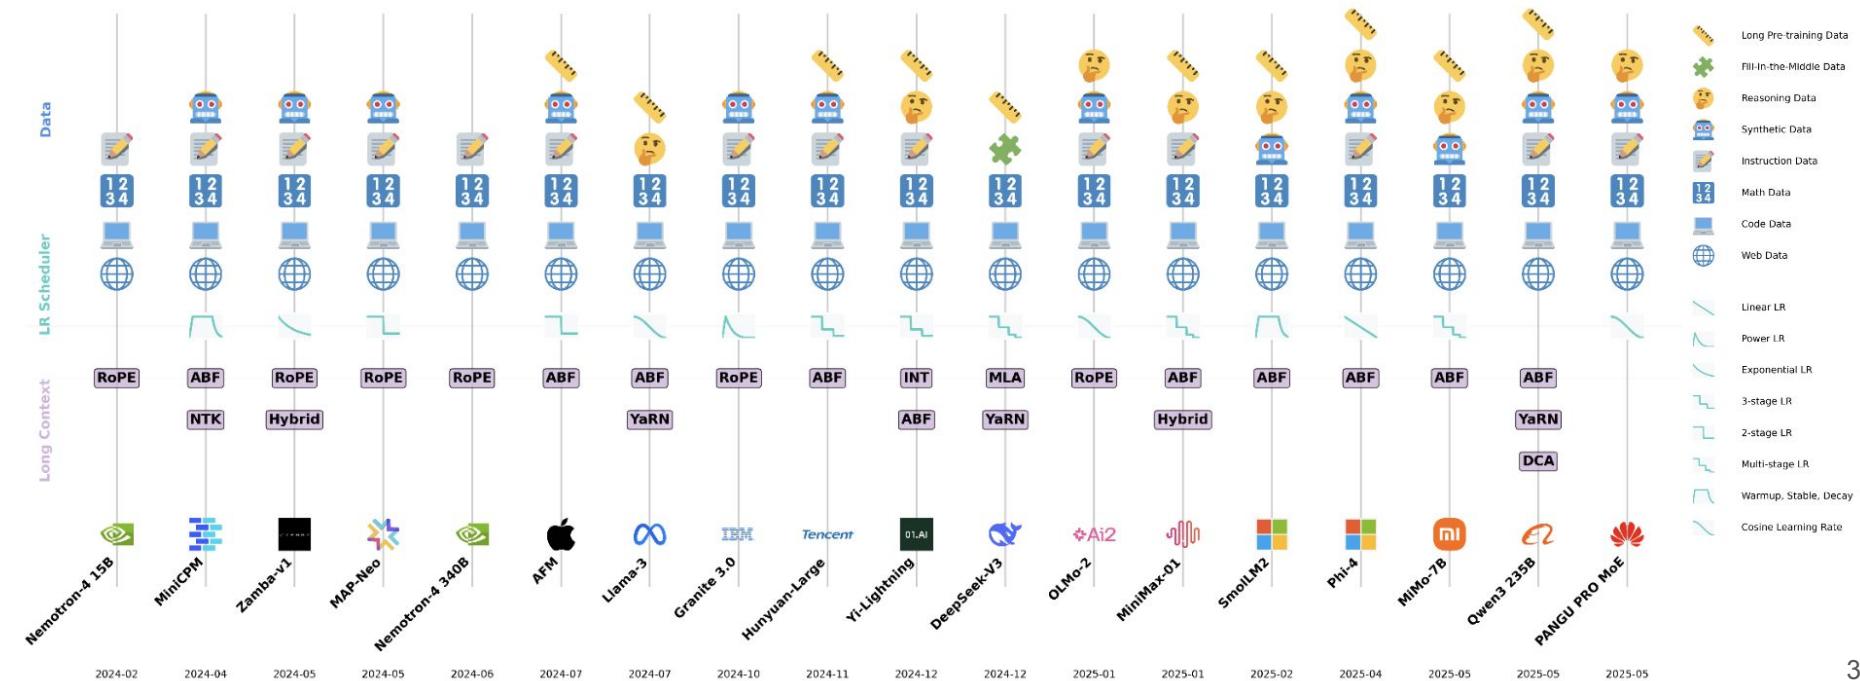
1. Pre-training

The step where the LLM learns to maximize the probability of the raw training data (*from the web*) is called ***pre-training***.



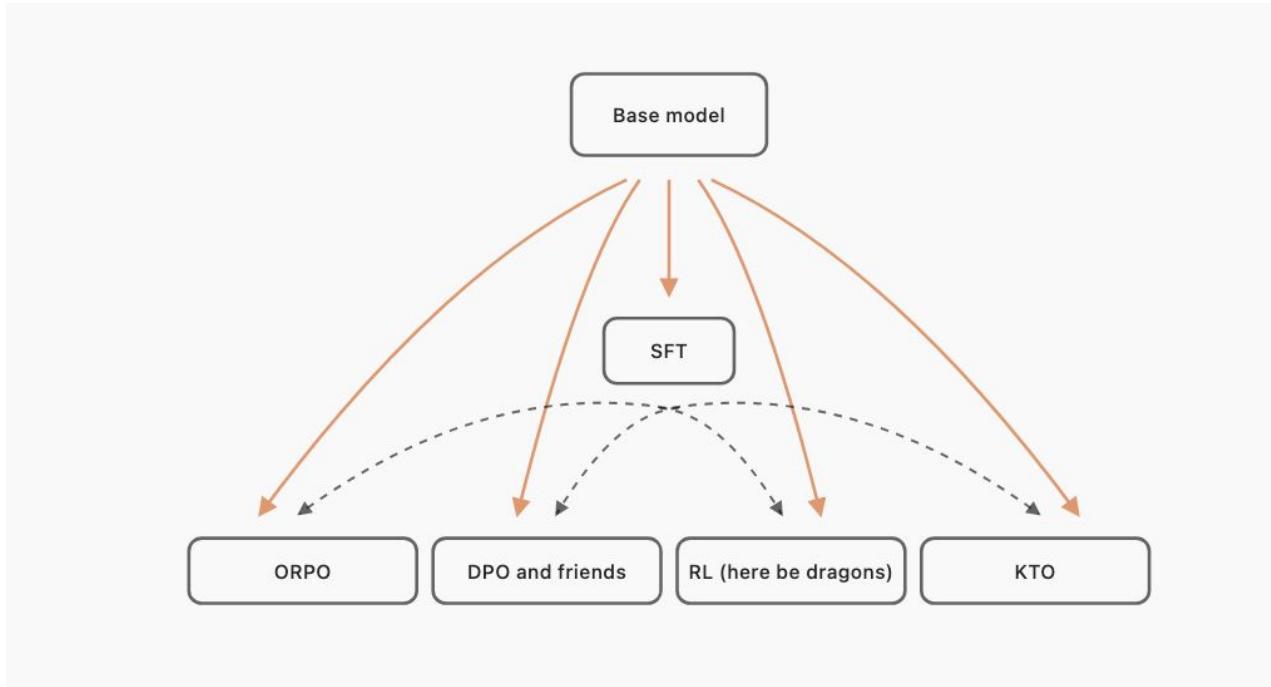
The different steps of LLM Training

2. Mid-training: Long-context extension, Multilingual data, annealing, etc.



The different steps of LLM Training

3. Post-training: Instruction following, preference alignment, agents, etc.



Sampling from an LLM

To fully benefit from the knowledge of the pre-trained LLM, we need to find good methods to sample tokens from the probability distribution.

This is not trivial.

You need to find a good trade-off between diversity of the generated texts, and coherence.

The general generation loop

```
def generate(prompt):
    generated ← prompt
    while True:
        next token  $t$  = sample ~  $p(\cdot | c=generated)$ 
        generated = generated +  $t$ 
        if  $t$  == end_of_sequence token:
            return generated
    continue
```

Greedy generation loop

```
def greedy(prompt):  
    generated ← prompt  
    while True:  
        # pick the token  $t$  with the highest score  
        next token  $t = \max \sim p(\cdot | c=generated)$   
        generated = generated +  $t$   
        if  $t ==$  end_of_sequence token:  
            return generated  
        continue
```

Top-K generation loop

```
def topk(prompt, k):
    generated ← prompt
    while True:
        # pick the token  $t$  with the highest score
        filtered_logits = keep the  $k$  highest scores
         $p(\cdot | c=generated)$  = softmax(filtered_logits)
        next token  $t$  = sample ~  $p(\cdot | c=generated)$ 
        generated = generated +  $t$ 
        if  $t$  == end_of_sequence token:
            return generated
    continue
```

For more diversity

Greedy generation algorithm will always output the same text

Topk will also very often output the same set of texts

Diversity is important for many reasons:

- creativity
- more exploration (we will see this for RL)
- avoid looping: “Hi, i am i am i am i am i am...”

The temperature hack

Observation: exponential in softmax “exaggerates” the differences in the logits

Solution: make the logits smaller by dividing them with a constant

```
logits = [10, 5, 1] # a list of logits of 3 tokens
```

without temperature

```
exp(logits) = [22026, 148, 2.7]
probabilities = [0.99, 0.01, 0.00]
```

with temperature $t = 5$

```
logits = [10, 5, 1] / 5 = [2.0, 1.0, 0.2]
exp(logits) = [7.39, 2.72, 1.22]
probabilities = [0.65, 0.24, 0.11]
# the token 2 and 3 have more chance to be sampled now!
```

Temperature generation loop

```
def temperature_generation(prompt, t):  
    generated ← prompt  
    while True:  
        # pick the token  $t$  with the highest score  
        temperature_logits = logits / t  
         $p(\cdot | \text{generated}) = \text{softmax}(\text{temperature\_logits})$   
        next token  $t = \text{sample} \sim p(\cdot | \text{generated})$   
        generated = generated +  $t$   
        if  $t == \text{end\_of\_sentence}$  token:  
            return generated  
        continue
```

Diversity / coherence trade-off

With previous sampling methods, the generation can drift quickly if we inadvertently sample the wrong token. We need a trade-off between diversity and coherence.

sorted_probs = [0.70, 0.10, 0.08, 0.05, 0.03, 0.02, 0.02]

Top-k (k=3): Samples from the first 3 tokens

- Problem: Includes token 3 with low score (0.08)
- If sampled, this weak token can drift the generation

Top-p (p=0.8): Samples from tokens where probability mass is concentrated

cumsum = [0.70, 0.80, 0.88, 0.93, 0.96, 0.98, 1.00]

Advantages:

- More adaptive: keeps only high-confidence tokens
- Prevents sampling from weak candidates

Top_p generation loop

```
def top_p(prompt, p):
    generated ← prompt
    while True:
        probabilites =  $p(\cdot | \text{generated})$ 
        cumsum_probabilites = cumsum(sorted(probabilities))
        # filter token that are <p
        logits = filter logits where cumsum_probabilites < p
         $p(\cdot | \text{generated}) = \text{softmax}(\text{logits})$ 
        next token  $t = \text{sample} \sim p(\cdot | \text{generated})$ 
        generated = generated +  $t$ 
        if  $t == \text{end\_of\_sentence}$  token:
            return generated
        continue
```

Conclusion

- The LLM estimates the probability distribution underlying the training data (the web), thanks to ***Maximum Likelihood Estimation*** by learning to predict next tokens
- After training, we can sample from the learned probability distribution
- Sampling is a trade-off between diversity and coherence:
 - **Greedy**: no diversity
 - **Temperature**: introduces diversity, but high temperature leads to gibish texts
 - **TopK**: coherence but less diversity, and the generation can drift
 - **TopP**: Adaptive to the probability distribution. A good trade-off between diversity and coherence. But it is not easy to find a good value for top_p.
 - **Many other sampling methods**: MinP, Typical Sampling, Beam-Search, etc.